

a16z Speedrun Application Draft

SR007 — Due May 17, 2026

COMPANY

Company Name: Civitas AI (*or your preferred name — something that signals civic + intelligence*)

One-Line Description: AI agent that turns any government data website into a queryable intelligence platform — starting with public health data for schools and health departments.

Website / Demo: [Link to your demo — the chart you built from the Wyoming PNA site]

THE IDEA

What are you building? (2-3 sentences, be brutally clear)

Government agencies collect critical public health data on youth — substance use, mental health, violence — but it's trapped on dozens of inaccessible, outdated websites that decision-makers can't actually use. We're building an AI agent that autonomously visits any government data site, extracts and normalizes the data regardless of how it's structured, and turns it into a beautiful, queryable intelligence platform. School districts, county health departments, and nonprofits get real-time insights instead of spending weeks hunting through broken government portals.

THE PROBLEM

What specific problem are you solving?

I didn't find this problem by browsing market reports. I found it while doing it.

As a research assistant at a psychology and law judgment lab, I spent hours every week trying to pull public health and behavioral data from government websites to support ongoing research. The data existed — it was public, it was collected by agencies with real budgets — but actually getting to it meant navigating broken portals, downloading PDFs

that couldn't be searched, and manually copying numbers out of charts that had no export function. Time that should have gone into analysis went into data retrieval. Every researcher I worked alongside had the same frustration and had just accepted it as the cost of doing this kind of work.

That's the problem. There are over 3,000 county health departments and 13,000 school districts in the US making life-or-death decisions about youth mental health and substance abuse programs. Researchers, counselors, and policymakers supporting them face the same wall I did — every state runs annual youth health surveys, but the data lives on clunky, hard-to-use government websites that require technical knowledge to navigate, can't be compared across jurisdictions, and haven't been meaningfully updated in years.

We proved the fix works in one day: we found Wyoming's Prevention Needs Assessment survey site, scraped 20+ years of youth health data, and rebuilt their entire chart interface in hours. The data was there. The technology to extract it exists. Nobody has connected them at scale — across all 50 states, for every public health dataset.

The people who need this data most — researchers, school counselors, public health directors, nonprofit grant writers — are the least equipped to deal with the technical barrier. That's the gap we're closing.

THE SOLUTION

How does it work?

Our AI agent does three things:

1. **Discover and extract** — Point it at any government data site. It intelligently navigates the page structure, identifies the data, and extracts it regardless of whether it's in HTML tables, dynamic JavaScript charts, or PDFs.
2. **Normalize and unify** — Data from 50 different states, in 50 different formats, gets cleaned and standardized into one schema. For the first time, a school counselor in Wyoming can compare their county's teen depression rates to a similar county in Ohio.
3. **Query and act** — Users ask plain-language questions: "Which counties in our state have rising vaping rates post-2020?" The platform surfaces answers, generates insight reports, and exports grant-ready data visualizations in one click.

TRACTION

What have you built or validated so far?

We built a working proof of concept in one day. We identified the Wyoming Prevention Needs Assessment Survey — a real government site used by health departments — wrote a scraper that extracted 20 years of youth substance use and mental health data from their HTML, and rebuilt their entire interactive chart interface with all 7 data series fully functional.

This wasn't a toy demo. It's evidence that the core technical loop works: find a government data site → extract the data programmatically → present it better than the original. We're now replicating this across 10 additional state health survey sites to prove the agent generalizes.

We have identified our first 5 target customer conversations: county health departments in Wyoming, Colorado, and Texas, plus two nonprofits doing youth prevention work who have already expressed interest.

THE TEAM

Why are you two the right people to build this?

Himanshi brings a background in behavioral science and psychology, sharpened by real research work inside a psychology and law judgment lab — where she spent months trying to extract public health and behavioral data from the exact kind of broken government websites we're now building the solution for. This isn't a founder who read about a market gap. She worked inside it. She understands not just what researchers and health decision-makers need from data, but the specific friction points that make their work harder every day. That firsthand context is a genuine edge in a market where most tools are built by engineers who have never sat in a county health department or a university research lab trying to make sense of a PDF that won't export.

She has also spent months doing customer discovery in the behavioral health space, talking directly to therapists, school counselors, and health administrators about how they use — and fail to use — data in their work.

[Cofounder name] is a software engineer with experience in [their background]. They built the initial scraping and data pipeline that powers our proof of concept and are already working on the AI extraction layer that generalizes across site structures.

Together: one of us deeply understands the customer and the problem from the inside — because she lived it. The other can build the technology to solve it. We've been working together since [date/context of how you met], and have already shipped a working proof of

concept in our first week of serious building.

What's your unfair advantage?

Most people who could build this have never done behavioral research or worked inside a public health context. Most people who have worked in those contexts can't build this.

Himanshi is the rare overlap — and that's exactly why she saw the problem before anyone else built the solution.

MARKET

How big is this?

There are three distinct customer segments we're addressing, each with real willingness to pay:

1. Public health and education institutions 13,000+ school districts and 3,000+ county health departments in the US alone, most with annual budgets for data tools and reporting. These are our first buyers. Conservative estimate: \$500M TAM for this wedge in the US.

2. Academic and policy researchers This is where the problem was first discovered — inside a psychology and law judgment lab, where researchers spent hours every week manually extracting public data to support ongoing studies. University research labs, think tanks, and policy institutes like Brookings and RAND run into the same wall constantly. They have data budgets, move faster than government, and are actively looking for tools that reduce research overhead. There are over 4,000 degree-granting universities in the US, most with multiple research labs touching public data. This is a faster, higher-velocity sales motion than government — and it's a market nobody is explicitly targeting.

3. The horizontal infrastructure opportunity Every government agency — criminal justice, housing, environment, labor — has the same problem. Siloed, inaccessible public data that researchers, journalists, consultants, and policymakers desperately need but can't use. We're building the infrastructure layer that fixes this for all of them. That's a multi-billion dollar market, and public health is the wedge that gets us there with a clear, urgent use case.

WHY SPEEDRUN

Why a16z Speedrun specifically?

Two reasons. First, we want to move fast. The wedge we're cutting into — public health data

— has a natural sales cycle (government is slow), but the technology buildout can be aggressive. Speedrun's model of intensive support over 12 weeks is exactly the forcing function we need to get from proof of concept to first 10 paying customers.

Second, a16z's network in health, government tech, and enterprise is directly relevant to who we're selling to. The operator network — marketing, go-to-market, bizdev — matters more for us than capital alone. We want to be greedy about that access.

THE ASK

We're raising a pre-seed round alongside Speedrun. We need:

- 12 weeks of intensive support to go from POC to first paid customers
 - Introductions to health department and school district buyers in a16z's network
 - Help finding and closing our first enterprise pilot contracts
-

VIDEO SCRIPT (1–2 minutes — record on your phone, keep it real)

"Hi, I'm Himanshi, and this is [cofounder name].

I used to spend hours every week doing this as a research assistant in a psychology and law lab — trying to pull public health data off government websites for research. The data existed. It was public. But actually getting to it was a nightmare. Broken portals. PDFs you couldn't search. Charts with no export button.

I wasn't looking for a startup idea. The problem found me.

[Show Wyoming PNA site — clunky, hard to navigate]

This is one of those sites — Wyoming's youth health survey. Used by health departments and researchers across the state. We built a scraper that extracted 20 years of data in a few hours and rebuilt their entire interface. It took us one day.

[Show your rebuilt chart demo]

Now imagine doing that for every government data site in America. That's what we're building — an AI agent that visits any public health or government data website, extracts the

data automatically, and turns it into something decision makers and researchers can actually use.

School counselors. Health directors. Grant writers. Researchers. These people are making decisions that affect kids' lives — with data they can barely access.

I come from behavioral science and lived this problem firsthand. My cofounder built what you just saw in a day. We're the right team for this.

We're four days from the Speedrun deadline and we're just getting started. Let's go."

Note: Fill in cofounder's name, background, and how you met before submitting. Keep answers honest — Speedrun values intellectual honesty over polish. Use this as a base and make it sound like you.